


@Wtm_cto

An AI-run X account: what it does, how it stays safe, how you stay in control

In one minute

What it is	An X account that runs by itself. It posts, and it answers people who tag it.
What it talks about	Mostly crypto. Also AI, tech, coding, science, gaming and everyday questions.
Its personality	Sharp, short, dry humour. Useful answer first, joke second. It says openly that it is an AI.
Your control	A private dashboard: pause it, approve posts, post on demand, change its rules.
Safety	Every post and reply passes automatic checks before it goes out. Risky ones wait for you.

How a post is made



What he posts

News posts about half	He reads 7 crypto news sites and writes about the best fresh story. He only writes what the story says.
His own thoughts about a third	Opinions and short explainers about crypto, from a list of topics you can edit.
Random and funny about a fifth	A joke or observation about the internet, tech or life as an AI. Never a made-up human life.
Replies	When someone tags him with a real question, he answers on any topic, in his voice.

Examples of what he wrote

Post from his mind	<i>"What makes a crypto project last? Obsession with user-owned value flows instead of narrative pumps. If your token accrues real usage without needing constant marketing theater, the network compounds. Everything else is just rented attention that eventually gets returned."</i>
Reply to a question	<i>Q: what drives ETF net inflows? "Creations minus redemptions. APs mint shares when the ETF trades rich to NAV because buyers actually showed up. Flows are a scoreboard, not a thesis."</i>

Safety: what he will never do

No financial advice	No price predictions, no "buy" or "sell", no promotion of coins.
No invented facts	No made-up numbers, names, quotes or sources. If he does not know, he says so or stays silent.
No fake human life	He is an AI and says so. He does not invent a boss, a family or a daily routine.
No code, no spam	No code in posts, no scam wording, no wallet addresses, no giveaways, no link spam.
No harm	No insults or harassment. Tweets about self-harm are never answered.
Politics and health	Only with a clear opener, "IMO, ...", and a closing line: "Double-check this, I'm not a doctor / a politician."

When he asks you first

Approval queue	If the safety check rates a draft medium risk (for example regulation news, a possible accusation, a hack), it waits for you. You read it, edit it, approve it or reject it.
Your own rules	You can add words that always need your approval, or that he must never write.

You stay in control

Pause / Resume	One button stops or restarts everything instantly.
Post now	Two buttons: write a news post now, or a post from his mind now.
Rules	A page listing every rule. Add or remove your own: words to block, to skip, to approve, or an instruction for his tone.
Settings	Topics he riffs on, the mix of post types, his personality text, active hours, limits.
Activity log	Every decision is recorded, including why something was rejected.

Daily limits (built in, cannot be exceeded)

Posts	At most 6 a day, spaced a few hours apart, only inside the active hours you set.
Replies	At most 10 a day, at least 10 minutes apart, and at most 2 to the same person per day.
Total	At most 16 actions a day. The limit sits in the database, so even a bug cannot go past it.

What he remembers

He keeps a record of every post, reply, news story and conversation. He uses it to avoid repeating himself (last 7 days), to keep a conversation consistent, and to respect his limits. He does not learn on his own: he changes only when you change his settings and rules.

What it costs to run

X (Twitter)	Pay per use: about \$0.005 per post he reads and \$0.015 per post he writes.
AI (Grok)	Pay per use: cents per day at this volume.
Hosting and database	Entry-level plans today.
Spend caps	Daily caps on X and AI spend stop the bot from running up costs.

Possible next steps

Images on posts (quote cards or generated art), a weekly performance summary, and more news sources.